



arm

Using Arm Development Tools to Enable Machine-Learning on Device ([blog](#))

Odin Shen | Staff FAE | Arm

NN on Arm?

7 2 1 0 4 1 4 9 5 9

- Deploy NN inference on Cortex v8.2 platform before real hardware?
- How to select solution?

Accuracy

Performance

Debug and Analysis

Availability

Capacity

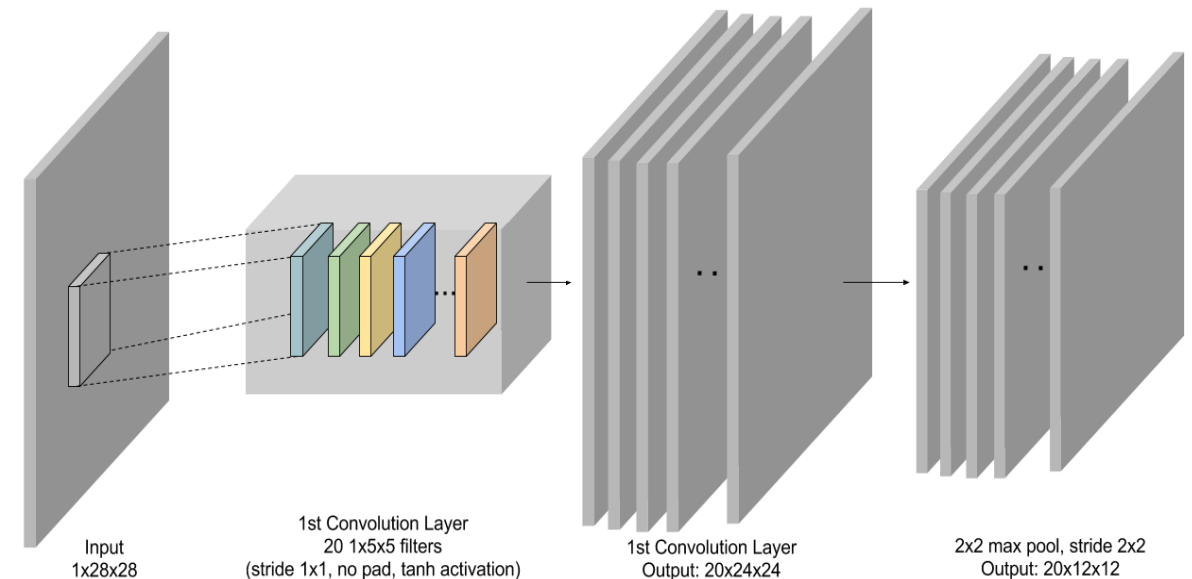

Cost

Flexibility

Use MNIST as Example

7 2 1 0 4 1 4 9 5 9

- The MNIST database of handwritten digits, available from this page, has a training set of 60,000 examples, and a test set of 10,000 examples.
- It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.


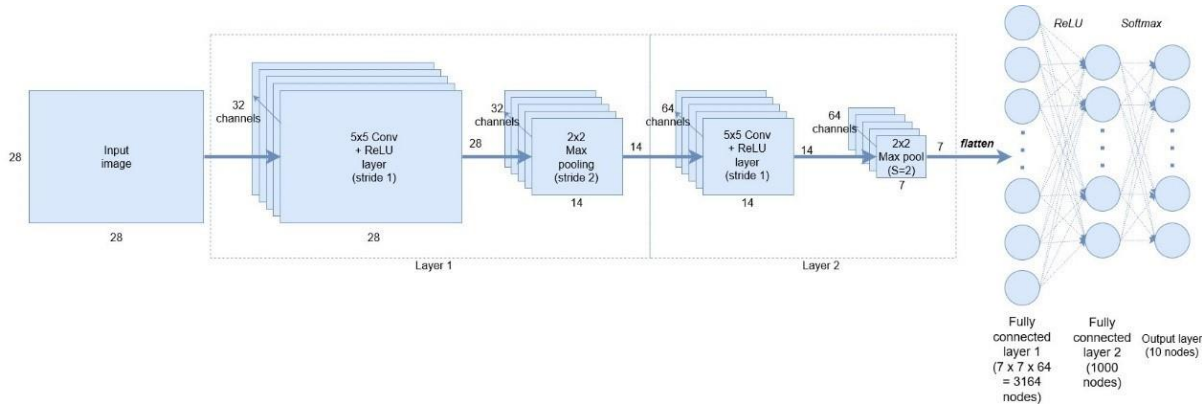


<http://yann.lecun.com/exdb/mnist/>

https://www.tensorflow.org/get_started/mnist/beginners

Use MNIST as Example

- Training on Tensor/Keras

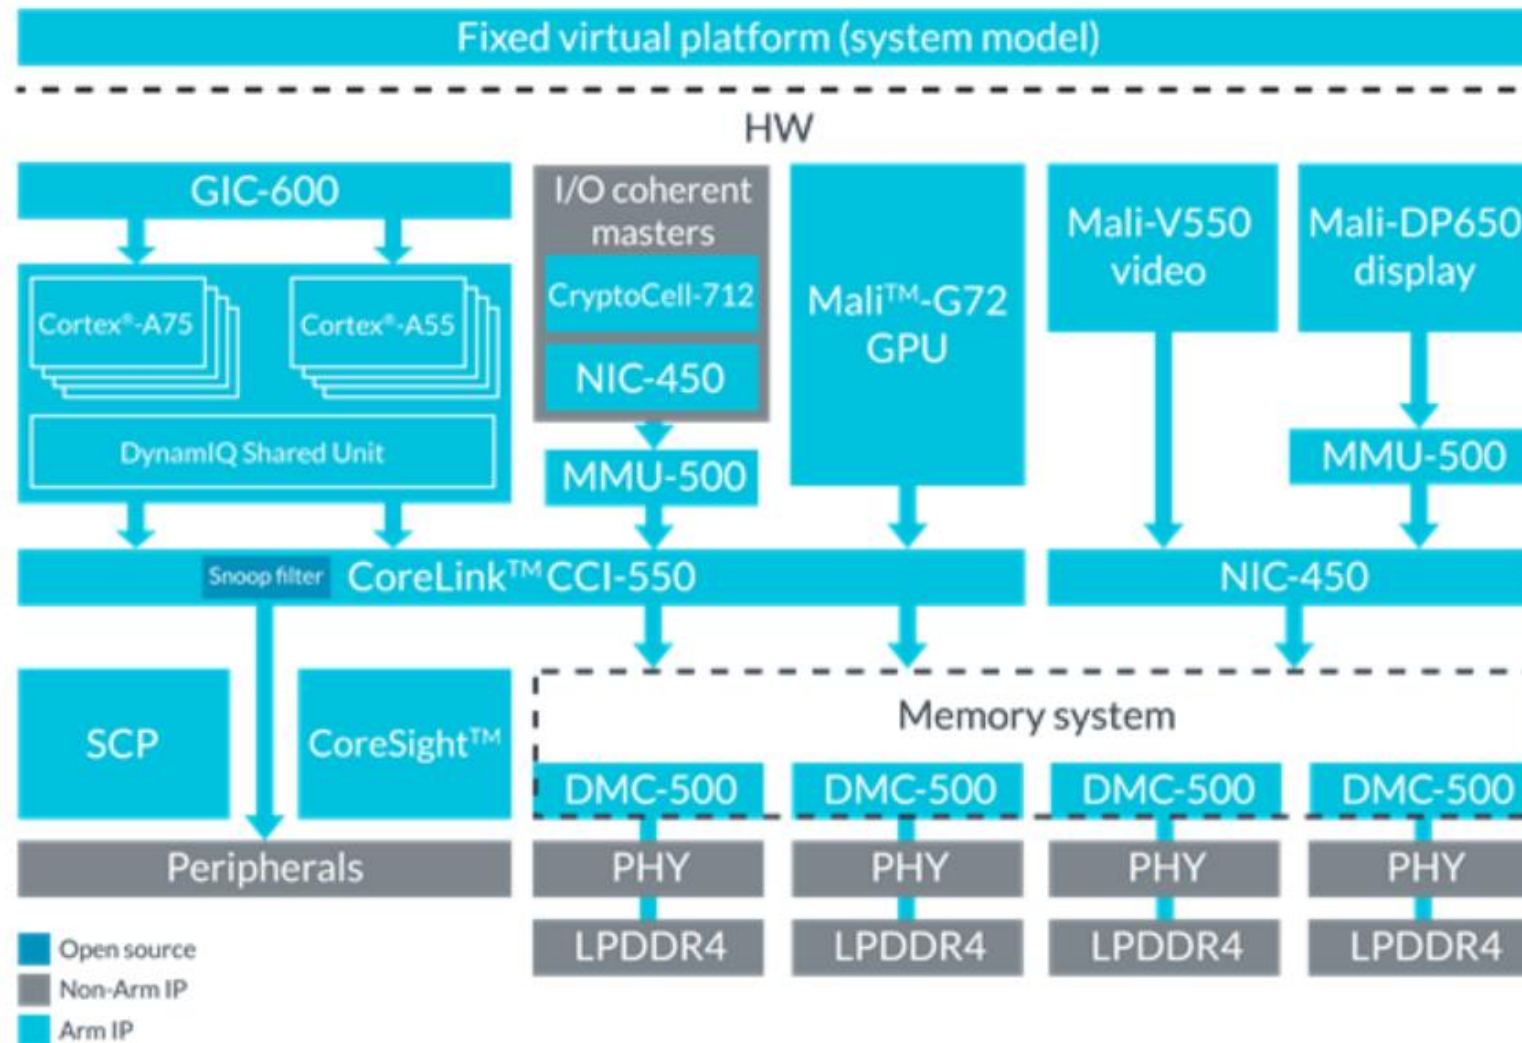


In [7]: `model.summary()`

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 24, 24, 16)	416
max_pooling2d_1 (MaxPooling2)	(None, 12, 12, 16)	0
conv2d_2 (Conv2D)	(None, 8, 8, 32)	12832
max_pooling2d_2 (MaxPooling2)	(None, 4, 4, 32)	0
dropout_1 (Dropout)	(None, 4, 4, 32)	0
flatten_1 (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 128)	65664
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290
Total params: 80,202		
Trainable params: 80,202		
Non-trainable params: 0		

Arm Fixed Virtual Platform (FVP)


System Guidance for Mobile 7210414959



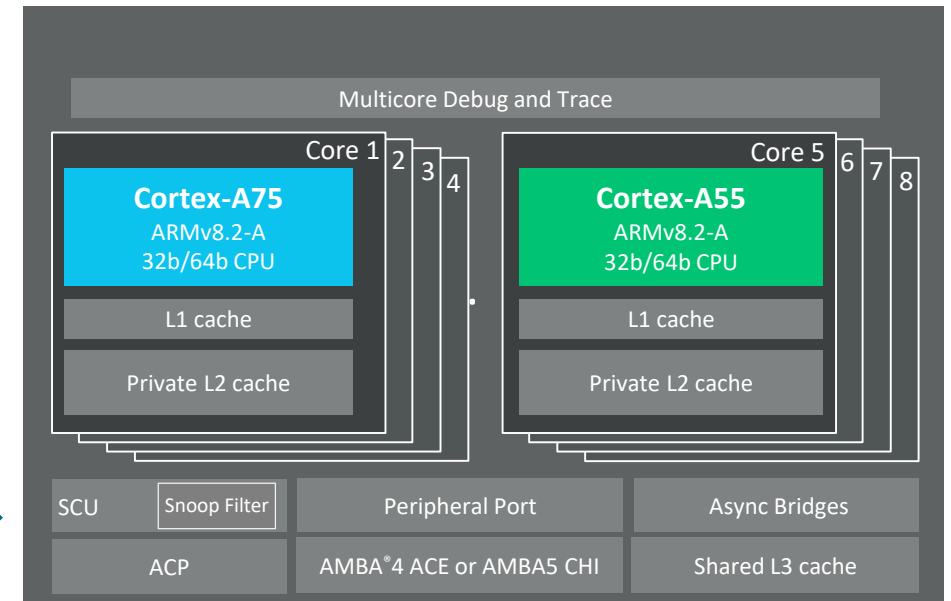
Early Deploy on Arm

7 2 1 0 4 1 4 9 5 9

- Purpose:
 - Prototyping on Arm platform
 - Early exploration NN engine design
 - SW framework profiling
- Script
 - Load image
 - Load parameter
- Automation
- Profiling



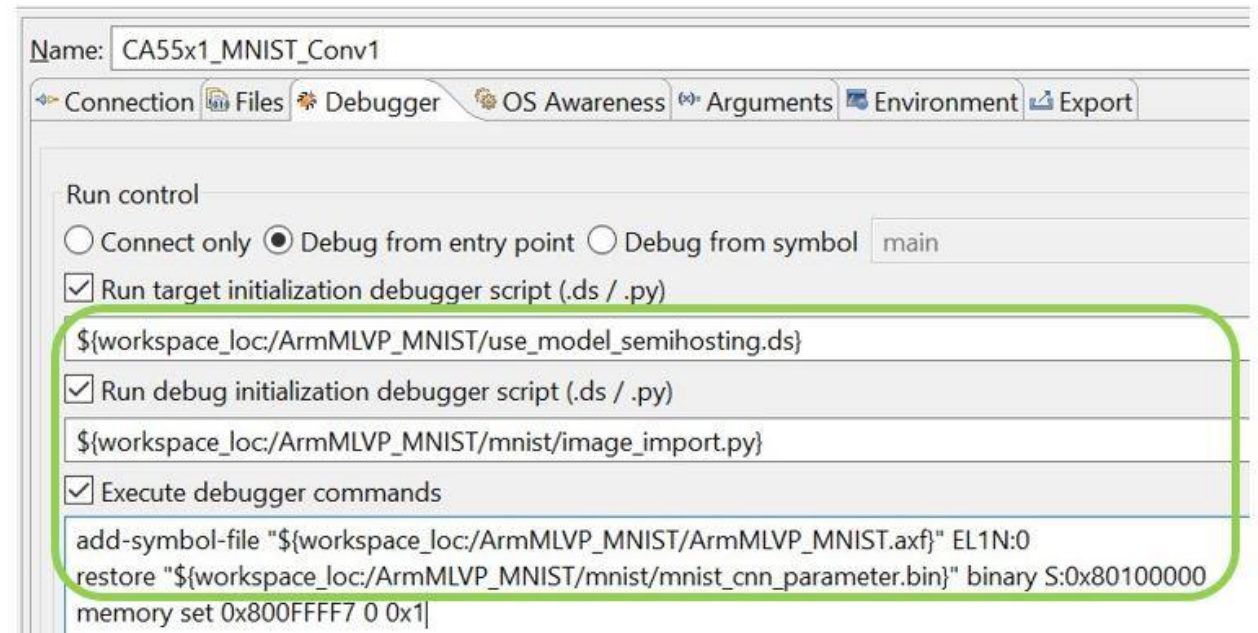
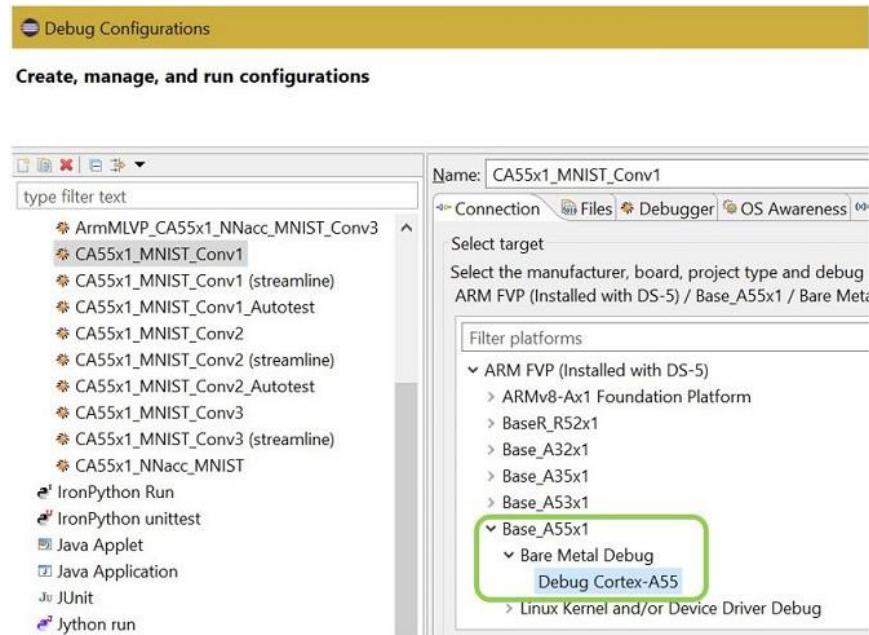
Debug
Trace
Streamline
Trace and Debug



Construct Platform

7 2 1 0 4 1 4 9 5 9

- Add a debug connection to hook on the bare-metal Cortex-A55 [Fixed Virtual Platforms](#) (FVP). You can treat virtual platform as a real silicon, all my following works are running on this.
- Select platform
- Scripting for image loading and H/W initialization



NN Performance Index

How to Analysis Your NN Device?

- Analysis your NN from four angles:
 - Execution time
 - NN code size,
 - CPU loading and
 - Memory access usage
- Memory access usage is another crucial factor of AI performance at the edge.

	Execution Time (instr)	NN code size (byte)	Convolution usage (%)	CPU Load/Store
#1				
#2				

Initial Version (O1)

Initial Version NN API Definition & Implementation

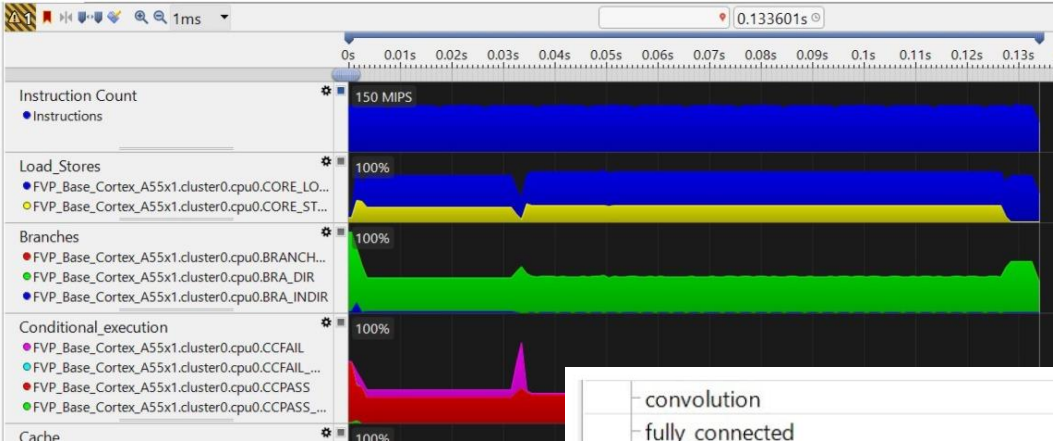
API	Description
convolution	Creates a convolution kernel that is convoluted with the layer input to produce a tensor of outputs
max_pooling	Reduce the number of parameters and amount of computation in the network
fully_connected	Connections to all activations in the previous layer, computed with a matrix multiplication followed by a bias offset

```
1 mnist_cnn_eval() {  
2     // Pre process  
3     convolution(&lay, layer0, layer1, layer0_paramter);  
4     max_pooling(&lay, layer1, layer2);  
5     convolution (&lay, layer2, layer3, layer2_paramter);  
6     max_pooling (&lay, layer3, layer4);  
7     fully_connected (&lay, layer4, layer5, layer5_paramter);  
8     fully_connected (&lay, layer5, layer6, layer6_paramter);  
9     // Post process  
10 }
```




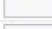





























NN#1 Result

- Test 

```
1 DS5 debug console
2 -----
3 Inf selected image [8] from CPU: 0
4   prob [ -16, -30, -6, -7, -18, -8, -21, -24, 17, -20,], avg/std: -13/10
5   Conv_mode: 1   selected image [8] from CPU: 0, inference: 8,      [Pass]
6     Instr count is 13146336
7       Cnt 0 is 0
8       Cnt 1 is 109812
9       Cnt 2 is 3416674
10      Cnt 3 is 13146336
11      Cnt 4 is 13146336
12
```




- Nested loop consume 96.7% CPU loading

- convolution	96.71%		1,292	96.71%		1,292	96.71%	
- fully_connected	1.35%		18	1.35%		18	1.35%	
- _flsbuf	0.67%		9	0.67%		9	0.67%	
- InvalidateUDCaches	0.60%		8	0.60%		8	0.60%	
- pmu_counter_get_event_type	0.22%		3	0.22%		3	0.22%	
- btod_internal_mul	0.07%		1	0.07%		1	0.07%	
- initTimerInterrupt	0.07%		1	0.07%		1	0.07%	
- _printf_fp_dec_real	0.07%		1	0.07%		1	0.07%	
- _printf_int_common	0.07%		1	0.07%		1	0.07%	
- _printf_int_dec	0.07%		1	0.07%		1	0.07%	
- _writebuf	0.07%		1	0.07%		1	0.07%	



Compiler Optimization Version

Compiler Optimization NN (O3)




convolution	86.99%	642	86.99%	642	86.99%
el1IrqHandler	7.32%	54	7.32%	54	7.32%
max_pooling	2.30%	17	2.30%	17	2.30%
_flsbuf	1.36%	10	1.36%	10	1.36%
InvalidateUDCaches	1.08%	8	1.08%	8	1.08%
...	0.14%	1	0.14%	1	0.14%

	Execution Time (instr)	NN code size (byte)	Convolution function usage	CPU Load/Store
O1	13146K	49K	96.7%	1187 / 547
O3	7166K	31K	86.99%	608 / 224

Deep Look CPU Load/Store Utilization

- Streamline provide the average CPU load/store count during the each of sampling to help the user has roughly number on any selected period. Use that we can roughly calculate the total load/store on whole NN operation.



Load: $8.33 \text{ K/s} * 73\text{ms} = 608$

Store: $3.08 \text{ K/s} * 73\text{ms} = 224$

- By tracing the assemble code on DS5, I can find the better coding to ...

Disassembly

```
0x0000000080004AFC: SUBS    w4, w4, #0x8
0x0000000080004B00: FMUL    v2.4s, v1.4s, v2.4s
0x0000000080004B04: FMUL    v3.4s, v1.4s, v3.4s
0x0000000080004B08: FADD    v2.4s, v4.4s, v2.4s
0x0000000080004B0C: FADD    v3.4s, v5.4s, v3.4s
0x0000000080004B10: ADD     w14, w14, #0x8
0x0000000080004B14: STP     q2, q3, [x24]
0x0000000080004B18: B.NE    0x80004ae8
0x0000000080004B1C: LDR     w4, [sp, #0x5c]
0x0000000080004B20: CMP     w25, w17
0x0000000080004B24: B.NE    0x80004ba8
0x0000000080004B28: B       0x80004bd4
0x0000000080004B2C: ADD     w14, w5, w18
0x0000000080004B30: ADD     w4, w16, w18
0x0000000080004B34: LDR     s0, [x1, w14, uxtw #2]
0x0000000080004B38: MUL     w14, w10, w4
0x0000000080004B3C: LSL     x19, x15, #2
0x0000000080004B40: LDR     s1, [x3, w14, uxtw #2]
0x0000000080004B44: LDR     s2, [x2, x19]
0x0000000080004B48: CMP     w22, #0x2
0x0000000080004B4C: FMUL    s1, s0, s1
0x0000000080004B50: FADD    s1, s2, s1
0x0000000080004B54: STR     s1, [x2, x19]
0x0000000080004B58: B.LO    0x80004bd4
0x0000000080004B5C: CMP     w25, #0x7
0x0000000080004B60: B.LS    0x80004ba4
0x0000000080004B64: MUL     w14, w22, w18
0x0000000080004B68: CMN     w27, w26
```

API Optimization Version

Re-design API

API	Description
convolution	Creates a convolution kernel that is convoluted with the layer input to produce a tensor of outputs
convolution_conv2()	Redesign convolution layer, dispatch each 2D input element calculation into convolution_filter2()
convolution_filter2()	Filter kernel convolution implementation
max_pooling	Reduce the number of parameters and amount of computation in the network
fully_connected	Connections to all activations in the previous layer, computed with a matrix multiplication followed by a bias offset

```

1 //pseudo code
2 float int convolution_filter2() {
3     // Load parameter
4     for (current_filter_row = 0; current_filter_row < filter_rows; current_filter_row++)
5         for (current_filter_col = 0; current_filter_col < filter_cols; current_filter_col++)
6             for (in_ch = 0; in_ch < input_channel; in_ch++) {
7                 current_input = ((float*)inputs)[ ((stride_row + current_filter_row) *
8                                                         + ((stride_col + current_filter_col) *
9                                                         + in_ch];
10                for (out_ch = 0; out_ch < output_channel; out_ch++) {
11                    current_weight = ((float*)weights)[ (current_filter_row * filter_c
12                                                         + (current_filter_col * input_ch
13                                                         + (in_ch
14                                                         + out_ch];
15                    current_result = current_input * current_weight;
16                    ((float*)outputs)[ (stride_row * output_columns * output_channel
17                                                         + (stride_col * output_channel
18                                                         + out_ch]
19                    += current_result;
20                }
21            }
22        }
23    }
24    for (out_ch = 0; out_ch < output_channel; out_ch++) {
25        current_biase = ((float*)biases)[out_ch];
26        kernel_output_addr = (stride_row * output_columns * output_channel) + (stride_c
27        kernel_result = ((float*)outputs)[kernel_output_addr];
28        kernel_result += current_biase;
29        if (relu_activation) {
30            kernel_result = relu(kernel_result);
31        }
32        ((float*)outputs)[kernel_output_addr] = kernel_result;
33    }
34 }
35 }
36
37 int convolution_conv2() {
38     // Initial
39     // Preload parameter
40     // Pre-processing
41     for (stride_row = 0; stride_row < lay->output_rows; stride_row++) {
42         for (stride_col = 0; stride_col < lay->output_columns; stride_col++) {
43             convolution_filter2(stride_row, stride_col, ...);
44         }
45     }
46 }

```

NN#3 Result




	Execution Time (instr)	NN code size (byte)	Convolution function usage	CPU Load/Store
O1	13146K	49K	96.7%	1187 / 547
O3	7166K	31K	86.99%	608 / 224
O3 w/conv2	3952K	28K	77.16%	250 / 76

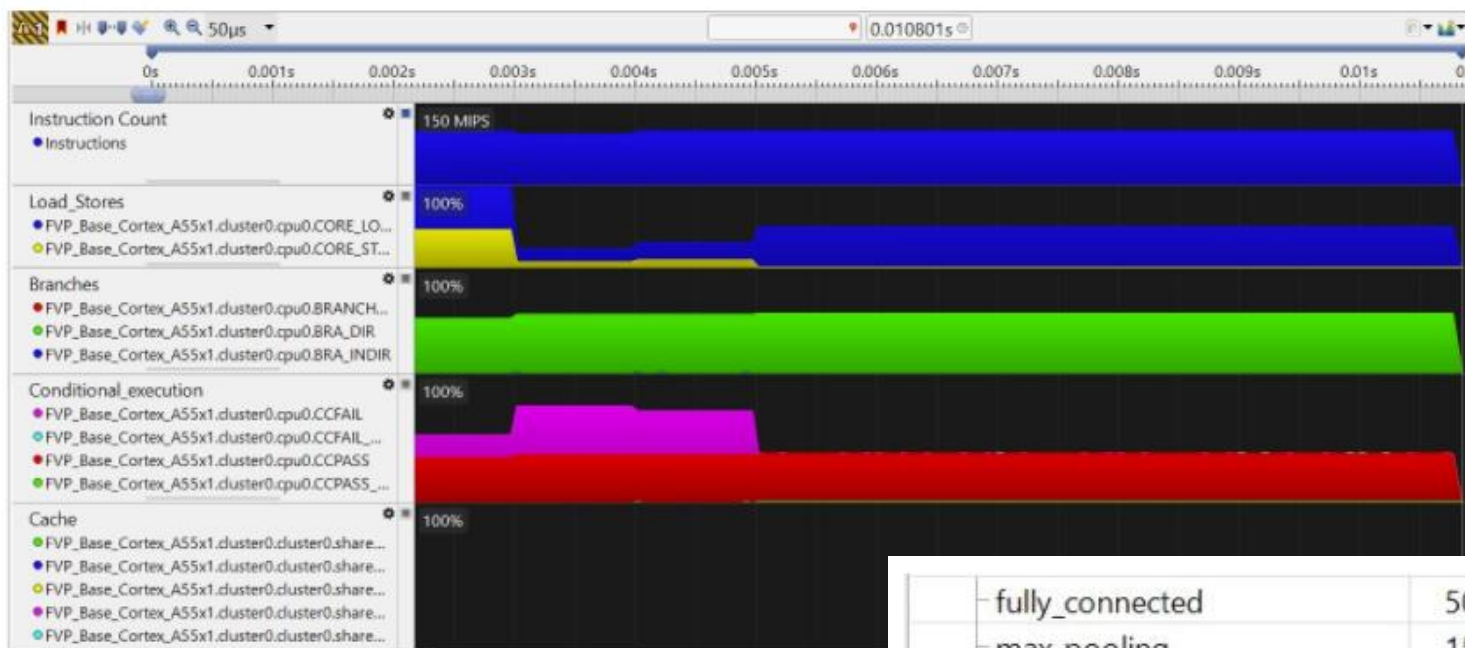
NN w/ Accelerator Model

Accelerator Concept

- One of straightforward concept is using NN accelerator to offload some of the work from the CPU, which has the capability to direct access input data, model parameter and write into the output buffer. Enable parallel 5x5 matrix multiplier.
- That's a straightforward idea, but the problem is I don't have RTL design to confirm with, does the same virtual platform environment can help with?
- Using SystemC/TLM to implement an idea of approximately timing behavior model and integrate into Fast Model.



NN#4 Result




fully_connected	50.00%	<div></div>	54	50.00%	<div></div>	54	50.00%	<div></div>
max_pooling	15.74%	<div></div>	17	15.74%	<div></div>	17	15.74%	<div></div>
convolution_conv3	12.96%	<div></div>	14	12.96%	<div></div>	14	12.96%	<div></div>
InvalidateUDCaches	7.41%	<div></div>	8	7.41%	<div></div>	8	7.41%	<div></div>
_flsbuf	6.48%	<div></div>	7	6.48%	<div></div>	7	6.48%	<div></div>
_printf_int_common	1.85%	<div></div>	2	1.85%	<div></div>	2	1.85%	<div></div>
puts	0.93%	<div></div>	1	0.93%	<div></div>	1	0.93%	<div></div>
_fputc\$unlocked	0.93%	<div></div>	1	0.93%	<div></div>	1	0.93%	<div></div>
_mutex_acquire	0.93%	<div></div>	1	0.93%	<div></div>	1	0.93%	<div></div>
_printf_f	0.93%	<div></div>	1	0.93%	<div></div>	1	0.93%	<div></div>
_printf_fp_dec_real	0.93%	<div></div>	1	0.93%	<div></div>	1	0.93%	<div></div>
_printf	0.93%	<div></div>	1	0.93%	<div></div>	1	0.93%	<div></div>

Summary

Final Result

	Execution Time (instr)	NN code size (byte)	Convolution function usage	CPU Load/Store
O1	13146K	49K	96.7%	1187 / 547
O3	7166K	31K	86.99%	608 / 224
O3 w/conv2	3952K	28K	77.16%	250 / 76
O3 w/NNacc	850k	24K	50%	106 / 23

CPU Loading Reduction Rate on different NN



Conclusion

- Using a sample NN application to demonstrate how to bring machine learning inference to an Arm device. This use case shows [DS-5](#) and [Fast Models](#) are excellent tools to help people develop and profile software algorithms on any Arm CPU.
- Welcome to visit [blog](#) and [code](#) for the detail.
- Next Step:
 - Adapt ML on Cortex-M by CMSIS-NN. (target MCU partner & ODM/OEM)
 - Compute Library profiling (target ISP partner)
 - ArmNN

Thank You!

Danke!

Merci!

谢谢!

ありがとう!

Gracias!

Kiitos!

감사합니다

धन्यवाद

arm